

rainfall-regression-1

September 15, 2023

1 Amount of Rainfall Prediction using Regression models

```
[1]: import pandas as pd
import numpy as np
import scipy.stats as stats
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: df = pd.read_csv("/content/rainfall in india.csv")
df.head()
```

```
[2]:
```

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	\
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	

	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	\
0	365.1	481.1	332.6	388.5	558.2	33.6	3373.2	136.3	560.3	
1	228.9	753.7	666.2	197.2	359.0	160.5	3520.7	159.8	458.3	
2	728.4	326.7	339.0	181.2	284.4	225.0	2957.4	156.7	236.1	
3	502.0	160.1	820.4	222.2	308.7	40.1	3079.6	24.1	506.9	
4	368.7	330.5	297.0	260.7	25.4	344.7	2566.7	1.3	309.7	

	Jun-Sep	Oct-Dec
0	1696.3	980.3
1	2185.9	716.7
2	1874.0	690.6
3	1977.6	571.0
4	1624.9	630.8

```
[3]: df.shape
```

```
[3]: (3887, 19)
```

```
[4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3887 entries, 0 to 3886
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   SUBDIVISION           3887 non-null   object
1   YEAR                  3887 non-null   int64
2   JAN                   3885 non-null   float64
3   FEB                   3885 non-null   float64
4   MAR                   3883 non-null   float64
5   APR                   3885 non-null   float64
6   MAY                   3886 non-null   float64
7   JUN                   3884 non-null   float64
8   JUL                   3883 non-null   float64
9   AUG                   3885 non-null   float64
10  SEP                   3884 non-null   float64
11  OCT                   3883 non-null   float64
12  NOV                   3882 non-null   float64
13  DEC                   3881 non-null   float64
14  ANNUAL                3872 non-null   float64
15  Jan-Feb               3884 non-null   float64
16  Mar-May               3882 non-null   float64
17  Jun-Sep               3881 non-null   float64
18  Oct-Dec               3880 non-null   float64
dtypes: float64(17), int64(1), object(1)
memory usage: 577.1+ KB

```

```
[5]: df.describe()
```

```

[5]:
count    YEAR    JAN    FEB    MAR    APR  \
count    3887.000000  3885.000000  3885.000000  3883.000000  3885.000000
mean     1958.221250   18.909833   22.165740   27.454365   41.072252
std        33.139584   33.790445   36.571112   47.818302   67.812487
min     1901.000000    0.000000    0.000000    0.000000    0.000000
25%     1930.000000    0.500000    0.500000    0.900000    2.800000
50%     1958.000000    5.700000    6.600000    7.300000   14.000000
75%     1987.000000   22.100000   27.300000   30.700000   45.500000
max     2015.000000  583.700000  403.500000  605.600000  595.100000

count    MAY    JUN    JUL    AUG    SEP  \
count    3886.000000  3884.000000  3883.000000  3885.000000  3884.000000
mean        79.227586  214.869902  338.604172  288.736371  196.910247
std       119.355452  226.145284  266.648228  189.613004  136.400297
min         0.000000   0.400000   2.400000   0.600000   0.100000
25%         7.900000  67.300000  171.450000  154.600000  100.375000
50%        33.000000  130.150000  279.000000  257.400000  172.900000
75%        83.850000  262.000000  407.750000  376.700000  265.400000

```

max	1168.600000	1609.900000	2362.800000	1664.600000	1222.000000
-----	-------------	-------------	-------------	-------------	-------------

	OCT	NOV	DEC	ANNUAL	Jan-Feb \
count	3883.000000	3882.000000	3881.000000	3872.000000	3884.000000
mean	87.588823	33.837816	17.057228	1361.243285	41.082878
std	93.031482	62.677156	41.071189	885.605666	60.263823
min	0.000000	0.000000	0.000000	62.300000	0.000000
25%	13.000000	0.500000	0.100000	785.025000	3.800000
50%	58.500000	8.000000	2.400000	1086.300000	18.950000
75%	135.700000	36.375000	14.700000	1543.225000	50.500000
max	948.300000	648.900000	617.500000	6331.100000	699.500000

	Mar-May	Jun-Sep	Oct-Dec
count	3882.000000	3881.000000	3880.000000
mean	147.419835	1038.634939	138.315541
std	200.307760	704.438118	153.514971
min	0.000000	57.400000	0.000000
25%	21.825000	558.700000	30.875000
50%	68.650000	861.200000	90.600000
75%	169.750000	1253.200000	190.950000
max	1745.800000	4536.900000	1252.500000

```
[6]: df.isnull().sum()
```

```
[6]: SUBDIVISION    0
YEAR              0
JAN               2
FEB               2
MAR               4
APR               2
MAY               1
JUN               3
JUL               4
AUG               2
SEP               3
OCT               4
NOV               5
DEC               6
ANNUAL            15
Jan-Feb           3
Mar-May           5
Jun-Sep           6
Oct-Dec           7
dtype: int64
```

```
[7]: df.fillna(method='bfill',inplace=True)
```

```
[8]: df.isnull().sum()
```

```
[8]: SUBDIVISION      0
YEAR                0
JAN                 0
FEB                 0
MAR                 0
APR                 0
MAY                 0
JUN                 0
JUL                 0
AUG                 0
SEP                 0
OCT                 0
NOV                 0
DEC                 0
ANNUAL              0
Jan-Feb             0
Mar-May             0
Jun-Sep             0
Oct-Dec             0
dtype: int64
```

```
[9]: df.duplicated().sum()
```

```
[9]: 0
```

```
[10]: df['SUBDIVISION'].value_counts()
```

```
[10]: EAST RAJASTHAN      115
WEST MADHYA PRADESH    115
NORTH INTERIOR KARNATAKA 115
COASTAL KARNATAKA      115
TAMIL NADU             115
RAYALSEEMA             115
TELANGANA              115
COASTAL ANDHRA PRADESH 115
CHHATTISGARH           115
VIDARBHA               115
MATATHWADA             115
MADHYA MAHARASHTRA     115
KONKAN & GOA           115
SAURASHTRA & KUTCH     115
GUJARAT REGION         115
EAST MADHYA PRADESH    115
SOUTH INTERIOR KARNATAKA 115
EAST UTTAR PRADESH     115
```

ASSAM & MEGHALAYA	115
JAMMU & KASHMIR	115
HIMACHAL PRADESH	115
PUNJAB	115
HARYANA DELHI & CHANDIGARH	115
UTTARAKHAND	115
WEST UTTAR PRADESH	115
WEST RAJASTHAN	115
BIHAR	115
JHARKHAND	115
ORISSA	115
GANGETIC WEST BENGAL	115
SUB HIMALAYAN WEST BENGAL & SIKKIM	115
NAGA MANI MIZO TRIPURA	115
ANDAMAN & NICOBAR ISLANDS	110
ARUNACHAL PRADESH	97

Name: SUBDIVISION, dtype: int64

```
[11]: df.YEAR.unique()
```

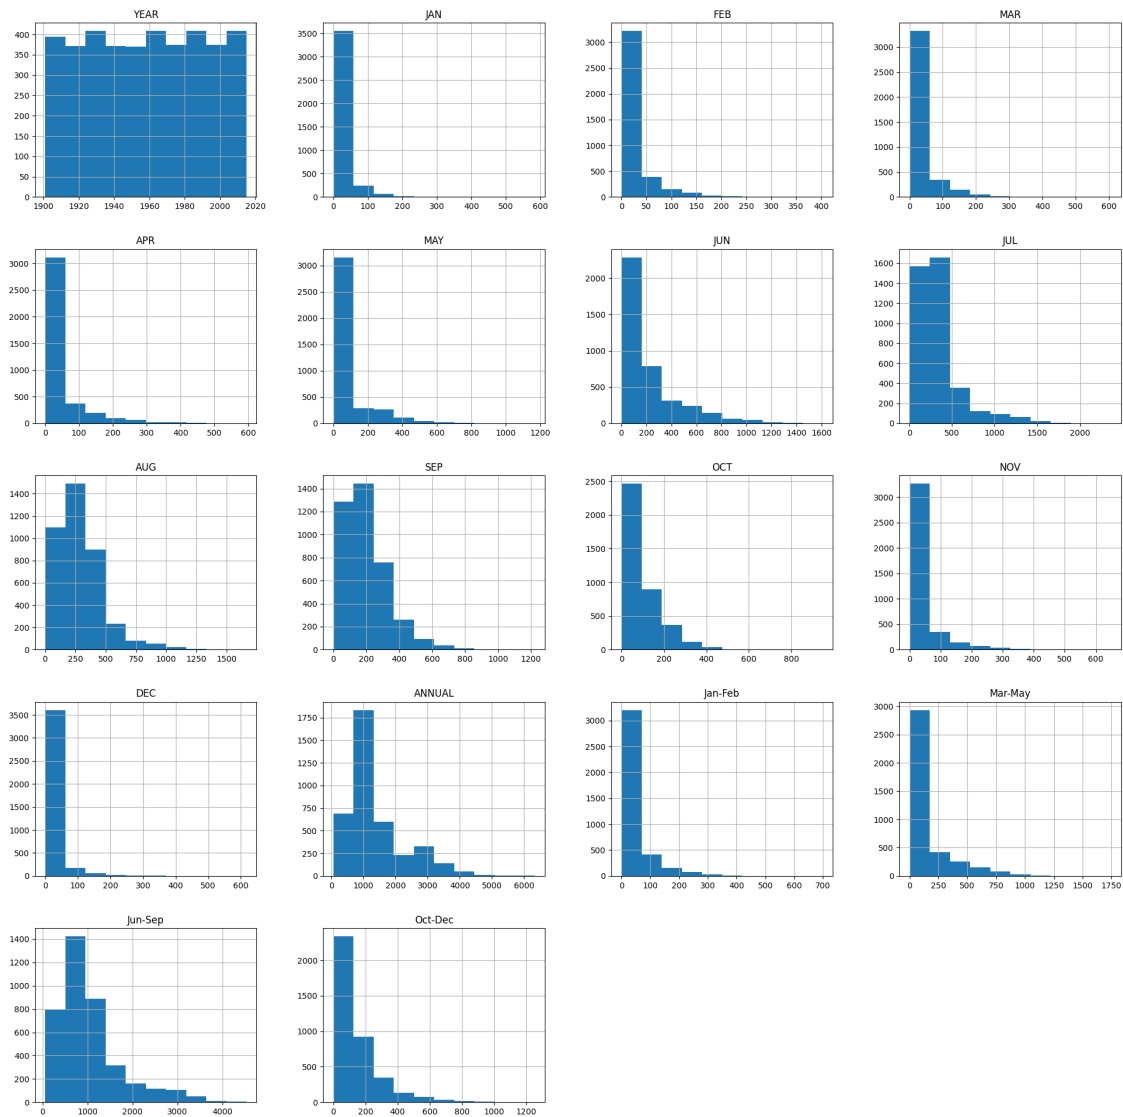
```
[11]: array([1901, 1902, 1903, 1904, 1905, 1906, 1907, 1908, 1910, 1911, 1912,
          1913, 1914, 1915, 1916, 1917, 1918, 1919, 1920, 1921, 1922, 1923,
          1924, 1925, 1926, 1927, 1928, 1929, 1930, 1931, 1932, 1933, 1934,
          1935, 1936, 1937, 1938, 1939, 1940, 1941, 1942, 1946, 1947, 1949,
          1950, 1951, 1952, 1953, 1954, 1955, 1956, 1957, 1958, 1959, 1960,
          1961, 1962, 1963, 1964, 1965, 1966, 1967, 1968, 1969, 1970, 1971,
          1972, 1973, 1974, 1975, 1976, 1977, 1978, 1979, 1980, 1981, 1982,
          1983, 1984, 1985, 1986, 1987, 1988, 1989, 1990, 1991, 1992, 1993,
          1994, 1995, 1996, 1997, 1998, 1999, 2000, 2001, 2002, 2003, 2004,
          2005, 2006, 2007, 2008, 2009, 2010, 2011, 2012, 2013, 2014, 2015,
          1943, 1944, 1945, 1948, 1909])
```

```
[12]: df.SUBDIVISION.unique()
```


```
[12]: array(['ANDAMAN & NICOBAR ISLANDS', 'ARUNACHAL PRADESH',
          'ASSAM & MEGHALAYA', 'NAGA MANI MIZO TRIPURA',
          'SUB HIMALAYAN WEST BENGAL & SIKKIM', 'GANGETIC WEST BENGAL',
          'ORISSA', 'JHARKHAND', 'BIHAR', 'EAST UTTAR PRADESH',
          'WEST UTTAR PRADESH', 'UTTARAKHAND', 'HARYANA DELHI & CHANDIGARH',
          'PUNJAB', 'HIMACHAL PRADESH', 'JAMMU & KASHMIR', 'WEST RAJASTHAN',
          'EAST RAJASTHAN', 'WEST MADHYA PRADESH', 'EAST MADHYA PRADESH',
          'GUJARAT REGION', 'SAURASHTRA & KUTCH', 'KONKAN & GOA',
          'MADHYA MAHARASHTRA', 'MATATHWADA', 'VIDARBHA', 'CHHATTISGARH',
          'COASTAL ANDHRA PRADESH', 'TELANGANA', 'RAYALSEEMA', 'TAMIL NADU',
          'COASTAL KARNATAKA', 'NORTH INTERIOR KARNATAKA',
          'SOUTH INTERIOR KARNATAKA'], dtype=object)
```

1.1 Visualizing the dataset

```
[13]: df.hist(figsize=(24,24));
```




```
[14]: df[["SUBDIVISION", "ANNUAL"]].groupby("SUBDIVISION").sum().
    ↪sort_values(by='ANNUAL', ascending=False).
    ↪plot(kind='barh', stacked=True, figsize=(15,10))
plt.xlabel("Rainfall in MM", size=12)
plt.ylabel("Sub-Division", size=12)
plt.title("Annual Rainfall v/s SubDivisions")
plt.grid(axis="x", linestyle="-.")
plt.show()
```



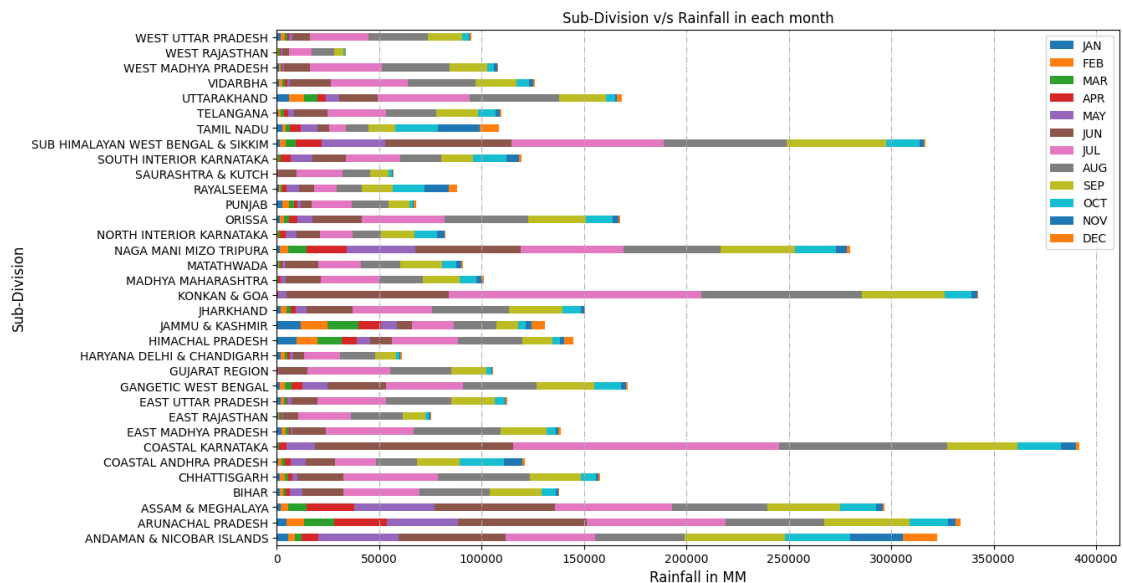
```
[15]: plt.figure(figsize=(15,8))
df.groupby("YEAR").sum()['ANNUAL'].plot(kind="line",color="r",marker=".")
plt.xlabel("YEARS",size=12)
plt.ylabel("RAINFALL IN MM",size=12)
plt.grid(axis="both",linestyle="-.")
plt.title("Rainfall over Years")
plt.show()
```

<ipython-input-15-64cf992d7485>:2: FutureWarning: The default value of numeric_only in DataFrameGroupBy.sum is deprecated. In a future version, numeric_only will default to False. Either specify numeric_only or select only columns which should be valid for the function.

```
df.groupby("YEAR").sum()['ANNUAL'].plot(kind="line",color="r",marker=".")
```



```
[16]: df[['SUBDIVISION', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL',
        'AUG', 'SEP', 'OCT', 'NOV', 'DEC']].groupby("SUBDIVISION").sum().
        plot(kind="barh",stacked=True,figsize=(13,8))
plt.title("Sub-Division v/s Rainfall in each month")
plt.xlabel("Rainfall in MM",size=12)
plt.ylabel("Sub-Division",size=12)
plt.grid(axis="x",linestyle="-.")
plt.show()
```




```
[17]: df.corr()
```

```
<ipython-input-17-2f6f6606aa2c>:1: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it will
default to False. Select only valid columns or specify the value of numeric_only
to silence this warning.
```

```
df.corr()
```

```
[17]:
```

	YEAR	JAN	FEB	MAR	APR	MAY	JUN	\
YEAR	1.000000	-0.049201	-0.019677	0.022538	0.009234	0.007683	-0.011065	
JAN	-0.049201	1.000000	0.466372	0.411304	0.222171	0.141304	-0.029891	
FEB	-0.019677	0.466372	1.000000	0.588000	0.381366	0.228551	0.050298	
MAR	0.022538	0.411304	0.588000	1.000000	0.571679	0.388697	0.174074	
APR	0.009234	0.222171	0.381366	0.571679	1.000000	0.677039	0.445185	
MAY	0.007683	0.141304	0.228551	0.388697	0.677039	1.000000	0.566982	
JUN	-0.011065	-0.029891	0.050298	0.174074	0.445185	0.566982	1.000000	
JUL	-0.013928	-0.043425	0.024944	0.094588	0.248373	0.324280	0.747678	
AUG	0.005350	0.021774	0.080199	0.135006	0.250687	0.344440	0.682612	
SEP	-0.011926	0.034833	0.084417	0.179542	0.390488	0.520697	0.587302	
OCT	0.007738	0.019887	0.010526	0.089758	0.351495	0.518810	0.443441	
NOV	-0.017265	0.078232	-0.009754	0.008894	0.128587	0.318825	0.143881	
DEC	-0.015863	0.239079	0.152872	0.147616	0.135111	0.250414	0.060478	
ANNUAL	-0.008958	0.123506	0.210254	0.337346	0.575650	0.700569	0.889232	
Jan-Feb	-0.039464	0.843778	0.868263	0.587386	0.355978	0.217881	0.013824	
Mar-May	0.012140	0.256634	0.407178	0.663610	0.877374	0.913959	0.529692	
Jun-Sep	-0.009532	-0.013208	0.063854	0.163197	0.380408	0.498240	0.901179	
Oct-Dec	-0.006275	0.107926	0.043385	0.097644	0.300292	0.509438	0.341960	

	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	\
YEAR	-0.013928	0.005350	-0.011926	0.007738	-0.017265	-0.015863	-0.008958	
JAN	-0.043425	0.021774	0.034833	0.019887	0.078232	0.239079	0.123506	
FEB	0.024944	0.080199	0.084417	0.010526	-0.009754	0.152872	0.210254	
MAR	0.094588	0.135006	0.179542	0.089758	0.008894	0.147616	0.337346	
APR	0.248373	0.250687	0.390488	0.351495	0.128587	0.135111	0.575650	
MAY	0.324280	0.344440	0.520697	0.518810	0.318825	0.250414	0.700569	
JUN	0.747678	0.682612	0.587302	0.443441	0.143881	0.060478	0.889232	
JUL	1.000000	0.699319	0.521876	0.261734	-0.019678	-0.035075	0.806421	
AUG	0.699319	1.000000	0.507431	0.255850	0.001576	-0.001542	0.772839	
SEP	0.521876	0.507431	1.000000	0.421982	0.172828	0.128717	0.740307	
OCT	0.261734	0.255850	0.421982	1.000000	0.425745	0.273332	0.556727	
NOV	-0.019678	0.001576	0.172828	0.425745	1.000000	0.465834	0.247656	
DEC	-0.035075	-0.001542	0.128717	0.273332	0.465834	1.000000	0.197774	
ANNUAL	0.806421	0.772839	0.740307	0.556727	0.247656	0.197774	1.000000	
Jan-Feb	-0.009163	0.060993	0.070789	0.017543	0.037971	0.226819	0.196934	
Mar-May	0.300626	0.322570	0.484315	0.448405	0.234523	0.228795	0.693827	
Jun-Sep	0.907246	0.850627	0.716369	0.392235	0.072689	0.030930	0.941727	


Oct-Dec 0.140230 0.153604 0.359787 0.849799 0.787910 0.621750 0.489474

	Jan-Feb	Mar-May	Jun-Sep	Oct-Dec
YEAR	-0.039464	0.012140	-0.009532	-0.006275
JAN	0.843778	0.256634	-0.013208	0.107926
FEB	0.868263	0.407178	0.063854	0.043385
MAR	0.587386	0.663610	0.163197	0.097644
APR	0.355978	0.877374	0.380408	0.300292
MAY	0.217881	0.913959	0.498240	0.509438
JUN	0.013824	0.529692	0.901179	0.341960
JUL	-0.009163	0.300626	0.907246	0.140230
AUG	0.060993	0.322570	0.850627	0.153604
SEP	0.070789	0.484315	0.716369	0.359787
OCT	0.017543	0.448405	0.392235	0.849799
NOV	0.037971	0.234523	0.072689	0.787910
DEC	0.226819	0.228795	0.030930	0.621750
ANNUAL	0.196934	0.693827	0.941727	0.489474
Jan-Feb	1.000000	0.390941	0.031418	0.086855
Mar-May	0.390941	1.000000	0.464635	0.427005
Jun-Sep	0.031418	0.464635	1.000000	0.274116
Oct-Dec	0.086855	0.427005	0.274116	1.000000

```
[18]: plt.figure(figsize=(15,6))
sns.heatmap(df.corr(),annot=True)
plt.show()
```

<ipython-input-18-2b948c0c3eb5>:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sns.heatmap(df.corr(),annot=True)
```



1.2 Label Encoding:

```
[19]: from sklearn.preprocessing import LabelEncoder
Encoder=LabelEncoder()
df["SUBDIVISION"] = Encoder.fit_transform(df["SUBDIVISION"])
```

```
[20]: df.head()
```

```
[20]:
```


	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	\
0	0	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	
1	0	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	
2	0	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	
3	0	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	
4	0	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	

	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	Jun-Sep	Oct-Dec
0	332.6	388.5	558.2	33.6	3373.2	136.3	560.3	1696.3	980.3
1	666.2	197.2	359.0	160.5	3520.7	159.8	458.3	2185.9	716.7
2	339.0	181.2	284.4	225.0	2957.4	156.7	236.1	1874.0	690.6
3	820.4	222.2	308.7	40.1	3079.6	24.1	506.9	1977.6	571.0
4	297.0	260.7	25.4	344.7	2566.7	1.3	309.7	1624.9	630.8


1.3 Removing Outliers:

```
[21]: #Plotting Boxplot for numerical value to identify whether outliers present or not
↳not
for i in df:
    print('Column = ',i)
    sns.boxplot(df[i])
    plt.show()
```


Column = SUBDIVISION




Column = YEAR




Column = JAN




Column = FEB




Column = MAR




Column = APR




Column = MAY




Column = JUN




Column = JUL




Column = AUG




Column = SEP




Column = OCT




Column = NOV




Column = DEC




Column = ANNUAL




Column = Jan-Feb




Column = Mar-May



Column = Jun-Sep



Column = Oct-Dec



```
[22]: Q1=df['ANNUAL'].quantile(0.25)
      Q3=df['ANNUAL'].quantile(0.75)
      IQR=Q3-Q1
      UL=Q3+1.5*IQR
      LL=Q1-1.5*IQR
      df=df[(df['ANNUAL']<UL) & (df['ANNUAL']>LL)]
```

```
[23]: df.shape
```

```
[23]: (3421, 19)
```

```
[24]: from datetime import date
      df['NO OFYEAR'] = date.today().year - df['YEAR']
```

<ipython-input-24-c6f53b52a45c>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df['NO OFYEAR'] = date.today().year - df['YEAR']

```
[25]: df.head()
```

```
[25]:
```

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	\
4	0	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	
5	0	1906	36.6	0.0	0.0	0.0	556.1	733.3	247.7	320.5	
9	0	1911	0.0	8.4	0.0	122.5	327.3	649.0	253.0	187.1	
11	0	1913	84.8	0.5	1.3	2.5	190.7	530.0	280.8	205.8	
15	0	1917	8.0	3.6	112.0	4.5	295.9	301.1	394.8	437.4	

	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	Jun-Sep	Oct-Dec	\
4	297.0	260.7	25.4	344.7	2566.7	1.3	309.7	1624.9	630.8	
5	164.3	267.8	128.9	79.2	2534.4	36.6	556.1	1465.8	475.9	
9	464.5	333.8	94.5	247.1	2687.2	8.4	449.8	1553.6	675.4	
11	580.1	288.8	133.0	67.5	2365.8	85.3	194.5	1596.7	489.3	
15	471.8	238.1	108.3	236.9	2612.4	11.6	412.4	1605.1	583.3	

	NO OFYEAR
4	118
5	117
9	112
11	110
15	106

1.4 Splitting the Data for Training Testing.

```
[26]: x=df.drop(columns=['ANNUAL','YEAR'])
      y=df['ANNUAL']
```

```
[27]: x.head()
```

```
[27]:
```

	SUBDIVISION	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	\
4	0	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	
5	0	36.6	0.0	0.0	0.0	556.1	733.3	247.7	320.5	164.3	
9	0	0.0	8.4	0.0	122.5	327.3	649.0	253.0	187.1	464.5	
11	0	84.8	0.5	1.3	2.5	190.7	530.0	280.8	205.8	580.1	
15	0	8.0	3.6	112.0	4.5	295.9	301.1	394.8	437.4	471.8	

	OCT	NOV	DEC	Jan-Feb	Mar-May	Jun-Sep	Oct-Dec	NO OFYEAR
4	260.7	25.4	344.7	1.3	309.7	1624.9	630.8	118
5	267.8	128.9	79.2	36.6	556.1	1465.8	475.9	117
9	333.8	94.5	247.1	8.4	449.8	1553.6	675.4	112
11	288.8	133.0	67.5	85.3	194.5	1596.7	489.3	110
15	238.1	108.3	236.9	11.6	412.4	1605.1	583.3	106

```
[28]: from sklearn.model_selection import train_test_split
```

```
[29]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.
      ↪3,random_state=42)
```

```
[30]: print(x_train.shape)
      print(x_test.shape)
      print(y_train.shape)
      print(y_test.shape)
```

```
(2394, 18)
(1027, 18)
(2394,)
(1027,)
```

2 Visualizing the results

```
[31]: import plotly.express as px
```

```
[32]: fig=px.scatter(df,x=df['SUBDIVISION'],
      ↪y=df['ANNUAL'],trendline='ols',trendline_color_override='red')
      fig.show()
      results = px.get_trendline_results(fig)
      print(results)
```

```
                                px_fit_results
0  <statsmodels.regression.linear_model.Regressio...
```

```
[33]: results.px_fit_results.iloc[0].summary()
```

```
[33]:
```

Dep. Variable:	y	R-squared:	0.092
Model:	OLS	Adj. R-squared:	0.092
Method:	Least Squares	F-statistic:	347.2
Date:	Fri, 15 Sep 2023	Prob (F-statistic):	6.85e-74
Time:	19:20:38	Log-Likelihood:	-26196.
No. Observations:	3421	AIC:	5.240e+04
Df Residuals:	3419	BIC:	5.241e+04
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	1409.8534	18.542	76.037	0.000	1373.500	1446.207
x1	-17.3091	0.929	-18.635	0.000	-19.130	-15.488

Omnibus:	512.921	Durbin-Watson:	0.306
Prob(Omnibus):	0.000	Jarque-Bera (JB):	794.577
Skew:	1.050	Prob(JB):	2.88e-173
Kurtosis:	4.080	Cond. No.	42.3

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[34]: fig=px.scatter(df,x=df['YEAR'],  
    ↪y=df['ANNUAL'],trendline='ols',trendline_color_override='red')  
fig.show()  
results = px.get_trendline_results(fig)  
print(results)
```

```
px_fit_results  
0 <statsmodels.regression.linear_model.Regressio...
```

```
[35]: fig=px.scatter(df,x=df['AUG'],  
    ↪y=df['ANNUAL'],trendline='ols',trendline_color_override='red')  
fig.show()  
results = px.get_trendline_results(fig)  
print(results)
```

```
px_fit_results  
0 <statsmodels.regression.linear_model.Regressio...
```

2.1 Scaling the data

```
[36]: #Scaling the data  
from sklearn.preprocessing import RobustScaler
```

```
[37]: rs=RobustScaler()  
x_train_scl=rs.fit_transform(x_train)  
x_test_scl=rs.transform(x_test)
```

```
[38]: from sklearn.metrics import  
    ↪mean_squared_error,mean_absolute_percentage_error,r2_score,mean_absolute_error
```

2.2 Fitting Linear Regression

```
[39]: from sklearn.linear_model import LinearRegression  
LR = LinearRegression()  
LR_model = LR.fit(x_train_scl,y_train)  
LR_model
```

```
[39]: LinearRegression()
```

```
[40]: y_pred = LR_model.predict(x_test_scl)  
y_pred_train = LR_model.predict(x_train_scl)
```

```
[41]: from math import sqrt
```

2.3 Evaluation metrics

```
[42]: mse=np.sqrt(mean_squared_error(y_test,y_pred))
      print('MSE = ',mse)
      mape=mean_absolute_percentage_error(y_test,y_pred)
      print('MAPE = ',mape)
      mae=mean_absolute_error(y_test,y_pred)
      print('MAE = ',mae)
```

```
MSE = 17.282893269004088
MAPE = 0.0012377277897482222
MAE = 1.4026635120204174
```

```
[43]: LR_rmse=np.sqrt(mean_squared_error(y_test,y_pred))
      print('RMSE = ',LR_rmse)
```

```
RMSE = 17.282893269004088
```

```
[44]: error=sqrt(mse)
      error
```

```
[44]: 4.157269929774117
```

```
[45]: r2_score(y_train,y_pred_train)
```

```
[45]: 0.9999773369945448
```

```
[46]: r2_score(y_test,y_pred)
```

```
[46]: 0.9990470027134594
```

2.4 Fitting Decision Tree Regression

```
[47]: from sklearn.tree import DecisionTreeRegressor
      DT = DecisionTreeRegressor()
      DT_model = DT.fit(x_train_scl,y_train)
      DT_model
```

```
[47]: DecisionTreeRegressor()
```

```
[48]: y_pred = DT_model.predict(x_test_scl)
      y_pred_train = DT_model.predict(x_train_scl)
```

```
[49]: mse=np.sqrt(mean_squared_error(y_test,y_pred))
      print('MSE = ',mse)
      mape=mean_absolute_percentage_error(y_test,y_pred)
      print('MAPE = ',mape)
      mae=mean_absolute_error(y_test,y_pred)
```

```
print('MAE =',mae)
```

```
MSE = 99.91395889506  
MAPE = 0.063134346421302  
MAE = 65.82882181110028
```

```
[50]: error=sqrt(mse)  
      error
```

```
[50]: 9.995697018970713
```

```
[51]: r2_score(y_test,y_pred)
```

```
[51]: 0.9681498778010226
```

3 Regularization Techniques

###-Ridge Regression:

```
[52]: from sklearn.linear_model import Ridge  
      from sklearn.model_selection import GridSearchCV
```

3.1 Tuning

```
[53]: ridge=Ridge()  
      parameters={'alpha': [1e-15,1e-10,1e-8,1e-3,1e-2,5,10,20,30,35,40,45,50,55,100]}  
      ridge_regressor=GridSearchCV(ridge,parameters,scoring='neg_mean_squared_error',cv=5)  
      ridge_regressor.fit(x_train_scl,y_train)
```

```
[53]: GridSearchCV(cv=5, estimator=Ridge(),  
                  param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.001, 0.01, 5, 10, 20,  
                                         30, 35, 40, 45, 50, 55, 100]},  
                  scoring='neg_mean_squared_error')
```

```
[54]: print(ridge_regressor.best_params_)  
      print(ridge_regressor.best_score_)
```

```
{'alpha': 5}  
-14.789731607600398
```

```
[55]: print("Best Parameter for Ridge:",ridge_regressor.best_estimator_)
```

```
Best Parameter for Ridge: Ridge(alpha=5)
```

```
[56]: ridge=Ridge(alpha=100.0)
```

```
[57]: ridge_model=ridge.fit(x_train_scl,y_train)
```

```
[58]: ridge_model
```

```
[58]: Ridge(alpha=100.0)
```

```
[59]: y_pred = ridge_model.predict(x_test_scl)
y_pred_train = ridge_model.predict(x_train_scl)
```

```
[60]: mse=np.sqrt(mean_squared_error(y_test,y_pred))
print('MSE = ',mse)
mape=mean_absolute_percentage_error(y_test,y_pred)
print('MAPE = ',mape)
mae=mean_absolute_error(y_test,y_pred)
print('MAE = ',mae)
```

```
MSE = 20.2976939702703
MAPE = 0.011965520401746576
MAE = 9.32958037824982
```

```
[61]: error=sqrt(mse)
error
```

```
[61]: 4.505296213377129
```

```
[62]: r2_score(y_train,y_pred_train)
```

```
[62]: 0.999576292865514
```

```
[63]: r2_score(y_test,y_pred)
```

```
[63]: 0.9986855254898491
```

###-Lasso Regression:

```
[64]: from sklearn.linear_model import Lasso
from sklearn.model_selection import GridSearchCV
```

```
[65]: # create a lasso object
lasso = Lasso(max_iter=10000)
```

```
[66]: # check for best alpha value using GridSearch
parameter={'alpha':[1e-15,1e-10,1e-8,1e-3,1e-2,1,5,1e1,1e2,1e3,1e4,1e5,1e6,1e7]}
lasso_regressor=GridSearchCV(
    lasso,parameter,
    scoring='neg_mean_squared_error',
    cv=5
)
```

```
[67]: lasso_model=lasso_regressor.fit(x_train_scl,y_train)
```

```
[68]: lasso_model
```

```
[68]: GridSearchCV(cv=5, estimator=Lasso(max_iter=10000),
                  param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.001, 0.01, 1, 5, 10.0,
                                         100.0, 1000.0, 10000.0, 100000.0, 1000000.0,
                                         10000000.0]}),
                  scoring='neg_mean_squared_error')
```

3.2 Tuning

```
[69]: GridSearchCV(cv=5, estimator=Lasso(max_iter=100000),
                  param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.001, 0.01, 1, 5, 10.0,
                                         100.0, 1000.0, 10000.0, 100000.0, 1000000.0,
                                         10000000.0]}),
                  scoring='neg_mean_squared_error')
```

```
[69]: GridSearchCV(cv=5, estimator=Lasso(max_iter=100000),
                  param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.001, 0.01, 1, 5, 10.0,
                                         100.0, 1000.0, 10000.0, 100000.0, 1000000.0,
                                         10000000.0]}),
                  scoring='neg_mean_squared_error')
```

```
[70]: print("Best Parameter for Lasso:",lasso_regressor.best_estimator_)
```

Best Parameter for Lasso: Lasso(alpha=0.01, max_iter=10000)

```
[71]: lasso=Lasso(alpha=100.0,max_iter=100000)

# fit into the object
lasso.fit(x_train_scl,y_train)
```

```
[71]: Lasso(alpha=100.0, max_iter=100000)
```

```
[72]: y_pred = lasso.predict(x_test_scl)
      y_pred_train = lasso.predict(x_train_scl)
```

```
[73]: mse=np.sqrt(mean_squared_error(y_test,y_pred))
      print('MSE = ',mse)
      mape=mean_absolute_percentage_error(y_test,y_pred)
      print('MAPE = ',mape)
      mae=mean_absolute_error(y_test,y_pred)
      print('MAE = ',mae)
```

MSE = 184.76371805122238
MAPE = 0.205593393948917
MAE = 146.77004671754315

```
[74]: error=sqrt(mse)
      error
```

```
[74]: 13.592781836372655
```

```
[75]: r2_score(y_train,y_pred_train)
```

```
[75]: 0.8877056551602704
```

```
[76]: r2_score(y_test,y_pred)
```

```
[76]: 0.8910838818566653
```

3.3 Fitting Linear SVR

```
[77]: from sklearn.svm import LinearSVR
```

```
[78]: SVR_mod = LinearSVR()
      SVR_model = SVR_mod.fit(x_train_scl,y_train)
      SVR_model
```

```
[78]: LinearSVR()
```

```
[79]: SVR_y_pred = SVR_model.predict(x_test_scl)
      SVR_y_pred_train= SVR_model.predict(x_train_scl)
```

```
[80]: r2_score(y_test,SVR_y_pred)
```

```
[80]: 0.9988244456890908
```

```
[81]: r2_score(y_train,SVR_y_pred_train)
```

```
[81]: 0.9999309279515805
```

4 Fitting Neural Networks

###MLP Regressor

```
[82]: from sklearn.neural_network import MLPRegressor
```

```
[83]: NN = MLPRegressor()
      NN_model = NN.fit(x_train_scl,y_train)
      NN_model
```

```
/usr/local/lib/python3.10/dist-
packages/sklearn/neural_network/_multilayer_perceptron.py:686:
ConvergenceWarning:
```

Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

```
[83]: MLPRegressor()
```

```
[84]: NN_y_pred = NN_model.predict(x_test_scl)
      NN_y_pred_train = NN_model.predict(x_train_scl)
```

```
[85]: r2_score(y_test, NN_y_pred)
```

```
[85]: 0.6166845881455115
```

```
[86]: r2_score(y_train, NN_y_pred_train)
```

```
[86]: 0.5725644900202689
```

4.1 Fitting Random Forest Regression

```
[87]: from sklearn.ensemble import RandomForestRegressor
```

```
[88]: RF = RandomForestRegressor(max_depth=100, max_features='sqrt',
      ↪min_samples_leaf=4,
      min_samples_split=10, n_estimators=800)
```

```
[89]: RF_model=RF.fit(x_train_scl, y_train)
      RF_model
```

```
[89]: RandomForestRegressor(max_depth=100, max_features='sqrt', min_samples_leaf=4,
      min_samples_split=10, n_estimators=800)
```

```
[90]: RF_y_pred=RF_model.predict(x_test_scl)
      RF_y_pred_train=RF_model.predict(x_train_scl)
```

```
[91]: mse=np.sqrt(mean_squared_error(y_test,y_pred))
      print('MSE = ',mse)
      mape=mean_absolute_percentage_error(y_test,y_pred)
      print('MAPE = ',mape)
      mae=mean_absolute_error(y_test,y_pred)
      print('MAE = ',mae)
```

```
MSE = 184.76371805122238
MAPE = 0.205593393948917
MAE = 146.77004671754315
```

```
[92]: error=sqrt(mse)
      error
```

```
[92]: 13.592781836372655
```

```
[93]: r2_score(y_test,RF_y_pred)
```

```
[93]: 0.9809028947052615
```

```
[94]: r2_score(y_train,RF_y_pred_train)
```

```
[94]: 0.9919404483485026
```

```
[94]:
```

4.1.1 Tracking Model Performance

Model Name	R2 score
Decission Tree	96.15
Linear Regression	98.70
Ridge Regression	97.51
Lasso Regression	89.10
Support Vector Reg	94.90
Neural Network(MLP)	61.66
Random Forest Regression	97.07

R2 scores for each of the models when running on the train/test data split.

```
[94]:
```